

## Blurring

- To blur across pixels, define a mask:
  - Whose values are non-negative
  - Whose value is largest at the center pixel
  - Whose entries sum to one.

## Edge Detection

- To find the edges in an image, define a mask:
  - Whose value is largest at the center pixel
  - Whose entries sum to zero.

## Forward Mapping – BAD!

- Iterate over source image

Multiple source pixels  
can map to same  
destination pixel

Some destination pixels  
may not be covered

## Reverse Mapping – GOOD!

- Iterate over destination image
  - Must resample source
  - May oversample, but much simpler!

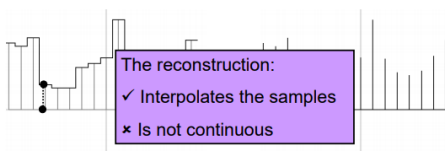
## Resampling

- Evaluate source image at arbitrary  $(u, v)$

$(u, v)$  does not usually  
have integer coordinates

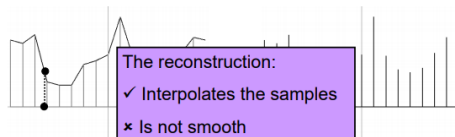
## Nearest Point Sampling

The value at a point is the value of the closest discrete sample.



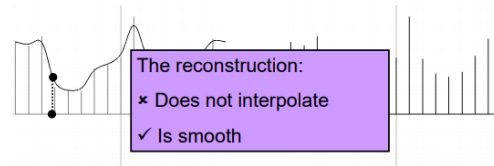
## Bilinear Sampling

The value at a point is the (bi)linear interpolation of the two surrounding samples.



## Gaussian Sampling


The value at a point is the Gaussian average of the surrounding samples.



## Image Sampling

Typically this is done in two steps:

1. Reconstruct a continuous function from input sample
2. Sample a continuous function at a fixed resolution.



Challenge:

Key Idea:  
Of all possible reconstructions, we want the one that is smoothest (has lowest frequencies).  
Signal processing helps us formulate this precisely.

## Question

- As higher frequency components are added to the approximation, finer details are captured.
- If we have  $n$  frequencies, how many samples can we fit?

Each frequency component has two degrees of freedom:  
• Amplitude  
• Shift  
With  $n$  frequencies we can fit  $2n$  samples.

## Sampling Theorem

Shannon's Theorem:

A signal can be reconstructed from its samples, if the original signal has no frequencies above  $1/2$  the sampling rate -- a.k.a. the *Nyquist Frequency*.

Terminology:

- A signal is *band-limited* if its highest non-zero frequency is bounded.
- The frequency is called the *bandwidth*.
- The minimum sampling rate for band-limited function is called the *Nyquist rate* (twice the bandwidth).

## Aliasing

- When a high-frequency signal is sampled with insufficiently many samples, it can be perceived as a lower-frequency signal. This masking of higher frequencies as lower ones is referred to as aliasing.

## Sampling

- There are two problems:
  - You don't have enough samples to correctly reconstruct your high-frequency information
  - You corrupt the low-frequency information because the high-frequencies mask themselves as lower ones.

## Summary – Nearest

- ✓ Can be implemented efficiently because the filter is non-zero in a very small region.
- ? Interpolates the samples.
- ✗ Is discontinuous.
- ✗ Does not address aliasing, giving bad results when a high-frequency signal is under-sampled.

## Summary – Linear

- ✓ Can be implemented efficiently because the filter is non-zero in a very small region.
- ? Interpolates the samples.
- ✗ Is not smooth.
- ✗ Partially addresses aliasing, but stills give bad results when a high-frequency signal is under-sampled.

## Summary – Gaussian


- ✗ Is slow to implement because the filter is non-zero in a large region.
- ? Does not interpolate the samples.
- ✓ Is smooth.
- ✓ Addresses aliasing by killing off high frequencies.

## Summary – Sinc

- ✗ Is slow to implement because the filter is non-zero in a large region.
- ? Does interpolate the samples.
- ✗ Assigns negative weights.
- ✗ Ringing at discontinuities.
- ✓ Addresses aliasing by killing off high frequencies.


## Alpha Channel

- Encodes pixel coverage information
  - $\alpha = 0$ : no coverage (or transparent)
  - $\alpha = 1$ : full coverage (or opaque)
  - $0 < \alpha < 1$ : partial coverage (or semi-transparent)
- Single Pixel Example:  $\alpha = 0.3$



Partial Coverage

or



Semi-Transparent

## Image Processing

- Quantization
  - Uniform Quantization
  - Random dither
  - Ordered dither
  - Floyd-Steinberg dither
- Pixel operations
  - Add random noise
  - Add luminance
  - Add contrast
  - Add saturation
- Filtering
  - Blur
  - Detect edges
- Warping
  - Scale
  - Rotate
  - Warp
- Combining
  - Composite
  - Morph


## Intersection Testing

Accelerated techniques try to leverage:

- Grouping: Discard groups of primitives that are guaranteed to be missed by the ray.
- Ordering: Test nearer intersections first and allow for early termination if there is a hit.


## Bounding Volume Hierarchies

- Use hierarchy to accelerate ray intersections
  - Intersect nodes only if you haven't hit anything closer



## Uniform (Voxel) Grid

- Construct uniform grid over the scene
  - Index primitives according to overlaps with grid cells



## Uniform (Voxel) Grid


- Trace rays through grid cells
  - Fast
  - Incremental

Only check primitives in intersected grid cells



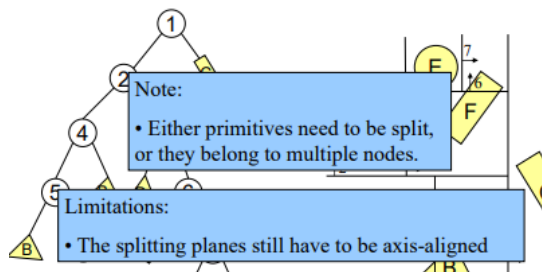
## Octrees

- In an octree, we only subdivide regions that contain more than one shape.
- Adaptively determines grid resolution.




## k-D Trees

- Alternate between splitting along the  $x$ -axis,  $y$ -axis, and  $z$ -axis.



## Binary Space Partition (BSP) Tree

- Example: Ray Intersection 2
  - Recursively split the ray and test nearer and farther halves, nearest first. Stop once you hit something:



## Transparency and Shadow

- Problem:
  - If a surface is transparent, then rays to the light source may pass through the object
  - Need to modify the shadow term so that instead of representing a binary (0/1) value, it gives the fraction of light passing through.
- ⇒ Accumulate transparency values as the ray travels to the light source.

## Snell's Law and Shadows

- Problem (Caustics):
  - If a surface is transparent, then rays to the light source may not travel in a straight line
  - This is difficult to address with ray-tracing

## Termination Criteria

- How do we determine when to stop recursing?
  - Depth of iteration
    - Bounds the number of times a ray will bounce around the scene
  - Cut-off value
    - Ignores contribution from bounces that contribute very little

## Termination Criteria

```
Pixel GetColor( scene , ray , ir , depth , cutOff )
{
    Pixel p(0,0,0)
    Ray reflect, refract
    Intersection hit = F
    if ( hit )
    {
        p += GetSu
        reflect.direction = Reflect( ray.direction , hit.normal )
        reflect.position = hit.position + reflect.direction * E
        if ( depth > 0 && hit.kSpec < cutOff )
            p += GetColor( scene , reflect , ir , depth-1 , cutOff/hit.kSpec ) * hit.kSpec
        refract.direction = Refract( ray.direction , hit.normal , ir , hit.ir )
        refract.position = hit.position + refract.direction * E
        if ( depth > 0 && hit.kTrans < cutOff )
            p += GetColor( scene , refract , ir , depth-1 , cutOff/hit.kTrans ) * hit.kTrans
    }
    return p
}
```

## Summary

- Ray casting (direct illumination)
  - Usually use simple analytic approximations for light source emission and surface reflectance
- Recursive ray tracing (global illumination)
  - Incorporate shadows, mirror reflections, and pure refractions

## Linear Transformations

- Linear transformations are combinations of ...
  - Scale, and
  - Rotation
- Properties of linear transformations:
  - Satisfies:  $T(s_1 \cdot p_1 + s_2 \cdot p_2) = s_1 \cdot T(p_1) + s_2 \cdot T(p_2)$
  - ⇒ Origin maps to origin
  - ⇒ Lines map to lines
  - ⇒ Parallel lines remain parallel
  - ⇒ Closed under composition

## Affine Transformations

- Affine transformations are combinations of ...
  - Linear transformations, and
  - Translations
- Properties of affine transformations:
  - Origin does not necessarily map to origin
  - Lines map to lines
  - Parallel lines remain parallel
  - Closed under composition

## Projective Transformations

- Projective transformations ...
  - Affine transformations, and
  - Projective warps

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- Properties of projective transformations:
  - Origin does not necessarily map to origin
  - Lines map to lines
  - Parallel lines do not necessarily remain parallel
  - Closed under composition

## Matrix Composition


- Be aware: order of transformations matters
  - » Matrix multiplication is not commutative

$$p' = \underbrace{T \cdot R}_{\text{"Global"}} \cdot \underbrace{S}_{\text{"Local"}} \cdot p$$


## Matrix Composition

- Rotate by  $\theta$  around arbitrary point  $(a, b)$ 
  - $M = T(a, b) \circ R(\theta) \circ T(-a, -b)$

The trick:  
First, translate  $(a, b)$  to the origin.  
Next, do the rotation about origin.  
Finally, translate back.



- Scale by  $(s_x, s_y)$  around arbitrary point  $(a, b)$ 
  - $M = T(a, b) \circ S(s_x, s_y) \circ T(-a, -b)$
  - (Use the same trick.)



## Scene Graphs

- Allow us to have multiple instances of a single model – reducing model storage size and making it easier to make consistent changes
- Allow us to model objects in local coordinates and then place them into a global frame – particularly important for animation
- Accelerate ray-tracing by providing a hierarchy that can be used for bounding volume testing

## Applying a Transformation

- Position

$$p' = M(p)$$

- Direction

$$\vec{v}' = M_L(\vec{v})$$

- Normal

$$\vec{n}' = (M_L^t)^{-1}(\vec{n})$$

$$\begin{matrix} \text{Affine} & & \text{Translate} & & \text{Linear} \\ \begin{pmatrix} a & b & c & t_x \\ d & e & f & t_y \\ g & h & i & t_z \\ 0 & 0 & 0 & 1 \end{pmatrix} & = & \begin{pmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{pmatrix} & \times & \begin{pmatrix} a & b & c & 0 \\ d & e & f & 0 \\ g & h & i & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} \\ M & & M_T & & M_L \end{matrix}$$

## Barycentric Coordinates

Barycentric coordinates are needed in:

- Ray-tracing, to test for intersection
- Rendering, to interpolate triangle information

## Barycentric Coordinates

Given the points  $p_1$ ,  $p_2$ , and  $p_3$ , how do we compute the barycentric coordinates of a point  $q$  in the plane spanned by  $p_1$ ,  $p_2$ , and  $p_3$ ?

(Signed) Area Ratios:

$$\alpha_q = \frac{A_1}{A_1 + A_2 + A_3}$$

$$\beta_q = \frac{A_2}{A_1 + A_2 + A_3}$$

$$\gamma_q = \frac{A_3}{A_1 + A_2 + A_3}$$

Solving this equation requires computing the areas of three triangles for every point  $q$ .

## Barycentric Coordinates

Given the points  $p_1$ ,  $p_2$ , and  $p_3$ , how do we compute the barycentric coordinates of a point  $q$  in the plane spanned by  $p_1$ ,  $p_2$ , and  $p_3$ ?

Solving this equation requires inverting a matrix.

However, the matrix is independent of the point  $q$  and can be computed once and re-used for all points  $q$ .

Ma

Recall:

The system can be singular even for non-degenerate triangles.

## Barycentric Coordinates

Linearity:

$$\begin{pmatrix} \beta \\ \gamma \end{pmatrix} = \begin{pmatrix} a_1 & a_2 & a_3 \\ b_1 & b_2 & b_3 \end{pmatrix} \begin{pmatrix} q_x \\ q_y \\ q_z \end{pmatrix}$$


If we set:

$$v_\beta = \begin{pmatrix} a_1 \\ a_2 \\ a_3 \end{pmatrix} \text{ and } v_\gamma = \begin{pmatrix} b_1 \\ b_2 \\ b_3 \end{pmatrix}$$

the barycentric coordinates can be expressed as dot-products:

$$\beta = \langle q, v_\beta \rangle$$

$$\gamma = \langle q, v_\gamma \rangle$$




## Perspective vs. Parallel

- Perspective projection
  - ✓ Size varies inversely with distance - looks realistic
  - ✓ Angles are preserved on faces parallel to the view plane
  - ✗ Distance are not preserved
  - ✗ Only parallel lines that are parallel to the view plane remain parallel
- Parallel projection
  - ✓ Good for exact measurements
  - ✓ Parallel lines remain parallel
  - ✓ Angles and distance are preserved on faces parallel to the view plane
  - ✗ Less realistic looking



## Cohen-Sutherland Line Clipping

- Associate an outcode to each vertex
- If both outcodes are 0, line segment is inside
- If AND of outcodes not 0, line segment is outside
- Otherwise clip and test

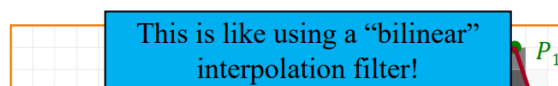


## Antialiasing Techniques

- Display at higher resolution
  - Corresponds to increasing sampling rate
  - Not always possible (fixed size monitors, fixed refresh rates, etc.)
- Modify pixel intensities
  - Vary pixel intensities along boundaries for antialiasing
  - Must have more than bi-level display

## Antialiasing

- Method 1: Area sampling
  - Calculate percent of pixel covered by primitive
  - Multiply this percentage by desired intensity/color



## Antialiasing

- Method 2: Supersampling (aka postfiltering)
  - Sample as if screen were higher resolution
  - Average multiple samples to get final intensity

## Antialiasing

Note that this makes things harder because pixels are no longer "owned" by a single triangle.

- Triangles contribute color rather than set color
- Along edges the total contribution must sum to one.
- ✗ Makes depth-testing more complicated.

## Flat Shading

- Can take advantage of spatial coherence
  - Make the lighting equation constant over the surface of each primitive


Surf	<ul style="list-style-type: none"> <li>• If the normal is constant over the primitive, and</li> <li>• if the light is directional, the diffuse component is the same for all points on the primitive</li> </ul>
Emissive	
Ambient	<ul style="list-style-type: none"> <li>• If the normal is constant over the primitive,</li> <li>• if the light is directional, and</li> <li>• if the direction to the viewer is constant over the primitive</li> </ul>
Diffuse	
Specular	<ul style="list-style-type: none"> <li>• If the normal is constant over the primitive, and</li> <li>• if the light is directional, and</li> <li>• if the direction to the viewer is constant over the primitive</li> </ul>

## Flat Shading

- Objects look like they are composed of polygons
  - OK for polyhedral objects
  - Not so good for smooth surfaces

## Gouraud Shading


- Produces smoothly shaded polygonal mesh
  - Continuous shading over adjacent polygons





## Phong Shading

- Linearly interpolate surface normals at vertices down and across scan lines



This was not supported in early generation graphic cards but can now be implemented in the fragment shader.

## Back-face detection

This method:

- Does not handle overlapping primitives
- Does not work for non-solid models and/or models without a well defined orientation.




## Ideal Solution

Painter's Algorithm:

- Sort primitives front to back and draw the back ones first, over-writing pixel values with information from the front primitives as they are processed.

Problem:


- In general you can't sort the primitives.
- ...Unless you are allowed to split them



## z-Buffer

- Store color & depth of closest object at each pixel
  - Initialize depth of each pixel to  $\infty$
  - Update only pixels whose depth is closer than in buffer
  - Depths are interpolated from vertices, just like colors

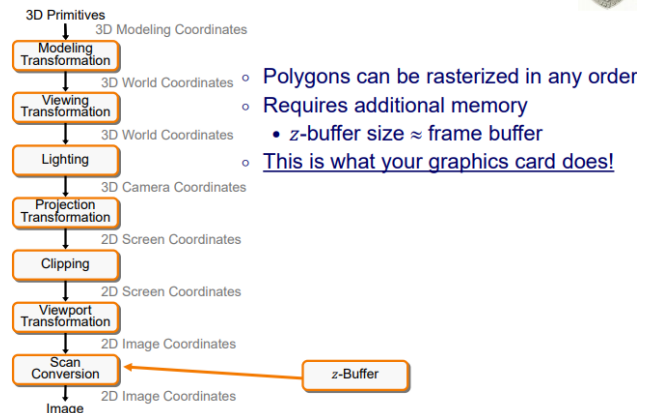
## Ray Casting Pipeline



Ray casting

- $P(p \log n)$  for  $p$  pixels and  $n$  shapes
- May (or not) use pixel coherence
- Simple, but generally not used

## 3D Rendering Pipeline



## Scan Conversion Example

A line segment in 2D projected onto a 1D screen.

- How do we interpolate correctly?

Recall: The 2D point  $(x, z)$  maps to the point  $(x/z)$  in 1D.

If  $p_1$  for a

To compute the interpolation weights correctly, we need to perform a perspective divide:

$$\frac{(1 - \alpha)x_1 + \alpha x_2}{(1 - \alpha)z_1 + \alpha z_2} = \frac{x}{1}$$

(1 -  $\alpha$ ) This is not the same as solving for the blending value in the image plane:

$$(1 - \alpha) \frac{x_1}{z_1} + \alpha \frac{x_2}{z_2} = \frac{x}{1}$$